

How Long Before the First Hit?

A conditioned-on-success estimate of time from education to a first \$10M-constant crossing, with its biases enumerated

Study repository: [~/trajectories](#)

23 August 2026

Abstract

We collected dated career anchors for 235 people drawn only from enumerated third-party lists, and measured the interval from the end of formal education to the first year an individual's own venture crossed \$10 million in **constant 2026 dollars**. Among the 93 people for whom that interval is computable on primary revenue evidence, the median is **9.0 years**, 95% bootstrap CI **[8.0, 10.0]**. Collection stopped under a pre-specified rule on median stability and interval width.

The estimate describes only people who eventually crossed. Those who worked for decades and never crossed are absent by construction, and the two strongest measured biases both push the figure short. **The number is a floor, not a forecast.** The finding we consider most useful is not the pooled median but its era split: careers whose hit landed before 1995 took a median of 13.0 years, those after 1995 took 8.0, and the sample is only 26% pre-1995.

*Reminder: every median in this paper is time-to-hit **among people who already succeeded**. The sample is conditioned on the outcome and has no denominator. No figure here supports any statement about the odds of succeeding.*

Reading the statistics in this paper

Four terms carry most of the weight. None of them mean what everyday usage suggests, so they are defined here in full.

Median. Line all 93 people up by how long they took and read off the person in the middle. Half took less, half took more. It is used here instead of the average because a few people took thirty years and an average would let them drag the figure upward; the median ignores how extreme the extremes are. *A median is a duration, never a probability.* “Median 9.0 years” means the middle person took nine years. It does not mean anyone has a nine-year timeline, and it says nothing about anyone's chance of getting there.

***n*.** How many people a figure is computed from. *n*=93 is the headline; where a table shows *n*=10, the figure beside it is an impression rather than a measurement, and this paper flags those explicitly.

95% confidence interval (CI). A range expressing how much the median would wobble if we had happened to sample different people. “[8.0, 10.0]” means the honest reading is *somewhere around 8 to 10 years*, not exactly 9.0. Loosely: if this study were repeated many times over, intervals built this way would contain the true value about 95% of the time. **A narrow interval means the number is well pinned down; a wide one means the**

data is thin. Whenever an interval appears, read the whole range — the single number in front of it is only its centre.

Half-width. Half the interval’s span, a one-number summary of how tight it is. [8.0, 10.0] spans 2.0, so the half-width is 1.0. This study’s stopping rule was written around it: collection continued until the half-width fell to 1.0 or below, i.e. until the estimate was pinned to roughly ± 1 year.

Bootstrap. How the interval is computed. Rather than assume the data follows a bell curve — it does not — the computer draws 10,000 random resamples of the 93 people, taking the median of each, and reads the interval off the spread of those 10,000 medians. It is a way of asking “how much would this number have moved if the sample had come out differently?” using only the data in hand.

IQR, or middle half. Where the interquartile range is quoted, it is the band containing the middle 50% of people, from the 25th to the 75th percentile. It describes *the spread of actual people*, and should not be confused with the confidence interval, which describes uncertainty about the median. The middle half of this sample crossed between 30 and 39 years of age; the interval on the median age is a much narrower [32.0, 34.0]. Both are true and they answer different questions.

The one that matters most

Every figure in this paper is computed over people who **already succeeded**. Statisticians call this conditioning on the outcome. It means the sample has no denominator: we know how long the winners took, and we have no idea how many people ran the same years and got nothing, because the study was built in a way that cannot see them.

So no number here is a rate, a chance, or a forecast. Nine years is how long it took the middle winner. It is not how long it takes, and it is not the odds.

1 What was measured

For each person we recorded six dated anchors, each with a source URL and a confidence grade: birth; the end of formal education (*last completed credential*, so a dropout is dated to their previous completed credential, not to the year they left); first paid work in the field the hit came from; first independent venture; **first hit**; and a much later scale milestone.

The headline clock is education-end to first hit. Two alternative clocks — from age 18, and from first venture — are reported in Section 3 because they survive different missing data.

1.1 The threshold is inflation-adjusted, and this matters more than it sounds

A hit is \$10 million in annual revenue **in constant 2026 dollars**. A nominal \$10M threshold is not the same test across eras: crossing \$10M in 1960 required a business roughly eleven times larger in real terms than crossing it today. The 1960 bar is **\$919,417**. Applied to a sample that is 26% pre-1995 by construction, a nominal threshold would have inflated exactly the careers this frame works hardest to include.

This was not a theoretical concern. In an early pilot a single row was dated eleven years late because a nominal \$10M was applied to a 1960s company. Every revenue year in the final data was checked against that year’s computed bar, and each row’s notes carry the pasted output of the conversion tool for the year used.

1.2 Selection

Names entered only by enumerating a named third-party list in its published order — rankings in rank order, prizes in year order — and taking entries as they came. Country, era, and sex were recorded but never used to select. The sole standard for entry is the money criterion.

2 Result

People researched	235
Included	201
Excluded	34 (14%)
Included on primary revenue evidence	121
Headline clock computable	93
Median education-to-hit	9.0 years
95% bootstrap CI	[8.0, 10.0]
CI half-width	1.00

Table 1: Headline result. The interval is a percentile bootstrap, 10,000 resamples, fixed seed.

Wave medians in collection order were 10.0, 9.0, 8.8, 9.0, 9.0, 9.8, 9.2, 9.0, 9.0. The estimate has been stable near nine years since roughly the fortieth person; the single excursion to 9.8 reversed on the following wave.

2.1 The era split is the more useful finding

Hit era	n	Median	95% CI
Pre-1995	33	13.0 yr	[9.0, 18.0]
Post-1995	60	8.0 yr	[6.5, 9.0]

Table 2: The pooled median is a mixture, not a constant.

The pooled 9.0 is not a stable property of ambitious careers. It is a weighted average of two eras that differ by five years, in a sample that is 74% post-1995. Anyone applying this number to a career today should read the 8.0, and anyone applying it historically should read the 13.0. Reporting the pooled figure alone would be the single most misleading thing this dataset permits.

We do not claim the mechanism. Faster capital formation, software distribution economics, and better documentation of recent private revenue would all produce this pattern, and the last of those is a measurement artifact rather than a change in the world.

3 Robustness

3.1 Missing education dates remove 23% of the eligible sample

Of 121 people with a hit dated on primary revenue evidence, **28 have no sourced education-end year** and therefore cannot enter the headline clock. This is the largest single threat to the estimate and it is structural: the headline clock is anchored to a credential, and a documented credential is exactly what is missing for self-taught and non-traditionally-educated founders.

The dropped rows are not concentrated in one region — 13 of 28 are US-based, against 48% of the retained rows — so the missingness does not appear to be a proxy for country. But it is

certainly not random with respect to education, and its direction on the median is unknown rather than benign.

Clock	n (of 121)	Median	95% CI
From education end	93 (77%)	9.0 yr	[8.0, 10.0]
From age 18	105 (87%)	15.0 yr	[14.0, 16.0]
From first venture	118 (98%)	4.2 yr	[4.0, 5.0]

Table 3: The two alternative clocks retain more of the sample. The venture clock, at 98% coverage and a half-width of 0.50, is the best-measured quantity in this study.

The venture clock is the one we would defend most strongly on coverage grounds: **once someone starts the thing that eventually works, the median time to \$10M-constant is 4.2 years**, and that figure rests on 98% of eligible rows rather than 77%.

3.2 The stopping rule fired on the boundary

Collection stopped under a rule fixed before the last waves: median drift below 0.50 years across the two most recent waves, and CI half-width at most 1.0 year. The final drift was (0.20, 0.00). The final half-width was **exactly 1.00**, against a test that fails only above 1.0.

Because that is uncomfortably close, we re-ran the bootstrap across 300 seeds. The rule passes in **272 of 300 (90.7%)**; the remaining 28 seeds give [8.0, 11.0], a half-width of 1.50, and would not have stopped. The seed was fixed in code in advance, so the stop is legitimate under the pre-specified procedure, but it is a boundary stop and we report it as one.

A second check separates the two things a boundary stop puts at risk. Adding one hypothetical row to the 93, at *any* clock value from 0 to 100 years, moves the median by **0.00 years** — a median over 93 observations simply does not care about one more. But a single added row of 12 years or more widens the interval to [8.0, 11.0], a half-width of 1.50, and the stopping rule would **not** have fired.

That is the honest characterisation of this study’s ending. **The estimate is robust and the stopping decision is not.** One more person — and the roster was exhausted, not the population — could have kept collection running without changing the number it would eventually report.

The natural criticism of a boundary stop is optional stopping — that collection ran until the number looked right. Against that, an interim report committed the day before collection ended, at 160 people and a revenue-strict n of 60, already reported a median of **9.2 years** ([analysis/archive/](#)). The estimate moved 0.2 years while the revenue-strict sample grew by more than half, and the wave-by-wave history shows the same stability from roughly the fortieth person. The stop determined *when* collection ended; it did not determine the number. A reader who prefers a stricter rule should treat collection as having ended arbitrarily at $n = 93$ and the interval as no tighter than [8.0, 11.0].

4 Bias catalogue

Twenty-four biases were catalogued during collection, several discovered only after they had corrupted data. The ones that bear on the number:

Survivorship, and it dominates everything else. The sample is conditioned on crossing. People who worked twenty years and never crossed are absent by construction, not by rarity.

This is not correctable within the design and is the reason no odds statement can be made.

Documentation bias, pushing short. A crossing is datable when revenue was published early. Companies that disclosed early skew toward faster growth and better funding. Undatable crossings became exclusions rather than late dates — 30 of the 34 exclusions are undatable crossings — so slow, quiet successes drop out rather than lengthening the tail.

Recall bias, corrected mid-study. An early roster agent recalled names and then found citations for them. Every row it produced had a valid citation and the sample was still biased, because recall came first. Later waves enumerate the list and read off it in published order.

Anti-fame steering, the same error inverted. Told to avoid names it would have recalled, a later agent rejected Peter Thiel, Vinod Khosla and Jensen Huang from enumerated lists as “too obvious.” That is the identical bias with its sign flipped. The rule is that prominence is not a criterion in either direction; the final wave took Phil Knight, Jensen Huang and John Hopfield in rank order and logged the temptation to skip them.

Correlated observations. Co-founders of one company share one crossing, and the bootstrap treats every row as independent, so a co-founder pair narrows the interval without adding information. A rule of one row per hit *event* — entity plus year — was added to the frame and rejected five candidates on its first use. One pair predating the rule remains, at Palm Computing in 1995; both rows are on fallback basis and therefore sit outside the headline set entirely, so they do not touch the reported interval.

Wave 1 was collected under a different design. Its composition was deliberately steered before the money-only rule was adopted. The sample is a mixture of two sampling designs and that discontinuity cannot be undone.

5 Audit method

Each wave was audited by re-researching a drawn sample of its rows **blind**: a second pass that could not see the first pass’s answers, run in a scratch directory containing only the brief and the inflation tool, after an early audit was found reading the first pass’s cached files in a shared directory.

Disagreement is scored as *contradiction* — two readings more than one year apart that do not overlap. A wave voids above 0.10 and its names are retired rather than reused.

Two rule changes were made during the study and both are disclosed here because each loosened the test. First, bounded ranges that overlap count as agreement: forcing a researcher to collapse an honest bracket to a single year to satisfy the check was voiding good waves. Second, a miss — one pass dating a row the other could not — stopped counting as disagreement. A six-platform blind cross-check measured a single-pass miss rate near 25%, so on ten audited rows two or three misses arise from search thoroughness alone; counting them voided waves that contained no wrong dates. Misses are now treated as rescue signals.

The final two audits returned contradiction rates of **0.000**.

6 What we would not defend

- **Annual recurring revenue may have been treated as calendar-year revenue.** One row is dated from ARR; a blind second pass using GAAP revenue gave a wider bracket that overlapped, so the audit did not flag it. Other SaaS rows may be dated the same way and were never swept.

- **Three rows rest on a single source** for the year that decides them, with the alternative reading disclosed in the row’s own notes.
- **Seven open questions were never resolved.** One concerns a founding year that, if corrected, would push a bounded range past the ten-year cap and turn that row into an exclusion. That row is outside the headline set, so the reported median does not move.
- **Thirty-five of the earliest roster rows** predate the rule requiring a membership URL naming the person.
- **Collection was executed by language models** following written briefs, with a human deciding methodology. The briefs, every batch log, every audit pairing and every rejection are in the repository.

7 Conclusion

Among people who eventually built something that crossed \$10 million in constant 2026 dollars, the median took **9.0 years from the end of formal education**, 95% CI [8.0, 10.0] — but 13.0 years if the hit landed before 1995 and 8.0 if after, and **4.2 years from starting the venture that worked**.

The most defensible reading is the narrowest one. This dataset measures how long the apprenticeship ran for people whose apprenticeship ended in a hit. It cannot say how many served one and got nothing, and the biases that survived enumeration push the surviving figure short. Treat nine years as a floor on the observed, not an expectation for the aspiring.